

## Unlocking the next frontier of personalized marketing

**As more consumers seek tailored online interactions, companies can turn to AI and generative AI to better scale their ability to personalize experiences.**

*This article is a collaborative effort by Eli Stein and Kelsey Robinson, with Alexis Wolfer, Gaelyn Almeida, and Willow Huang, representing views from McKinsey's Growth, Marketing & Sales Practice.*

**Here's the challenge facing brands and retailers:** Communicate clearly with a vast array of consumers who speak thousands of languages, hail from countless different cultures and socioeconomic backgrounds, and make purchasing decisions based on highly personal preferences. It's no easy feat to reach all of these consumers on a broad scale in an authentic way.

For some time now, companies have been trying to address customer needs through personalization, using data and analytics to craft more relevant consumer experiences. The goal is to present consumers with compelling offers and tailored, resonant messages at the right time. Today's customers want more of this. As previous McKinsey research revealed, 71 percent of consumers expected companies to deliver personalized interactions, and 76 percent got frustrated when it didn't happen.<sup>1</sup> When companies get it right, however, they can create significant value.

Companies often deploy tactical, manual, and stand-alone solutions to engage their customers. But retailers are now entering a promising new era of personalization. To reach consumers where they are and how they want to be met, marketers can embrace two powerful innovations: AI-driven targeted promotions, and the use of gen AI to create and scale highly relevant messages with bespoke tone, imagery, copy, and experiences at high volume and speed.

These innovations lay the groundwork for growth. Using improved analytics models, brands and retailers can better provide valuable offers to microcommunities wherever they want to engage. Meanwhile, gen AI enables marketers to create tailored content that

---

<sup>1</sup> "The value of getting personalization right—or wrong—is multiplying," McKinsey, November 12, 2021.

is relevant to those groups. Brands and retailers can better connect with customers by using language that speaks to them and by providing communications that resonate and give consumers a reason to engage.

To unlock the potential of targeted promotions and content, marketers should prioritize efforts to boost their underlying marketing technology stack. A robust framework built on better data, decisioning, design, distribution, and measurement is essential. With improved analysis through technology, marketers can gain deeper insights into customer behaviors and preferences, provide improved personalized experiences, and incorporate tactics that feed into a long-term personalization strategy for growth.

## The promise of targeted promotions

For both companies and customers, the old way of managing promotions—blunt offers to large groups of people—is no longer cutting it. Retailers face pressures due to economic uncertainty, changing consumer preferences, and, in some cases, declining profits. Meanwhile, previous McKinsey research suggests that 65 percent of customers see targeted promotions as a top reason to make a purchase.

Many retailers view AI and gen AI as a way to reverse the downward trends and accelerate growth. An increasing number are starting to experiment with AI to improve *mass* promotions. But companies can be more strategic by employing AI for *targeted* promotions, using data to tailor discounts based on people's shopping preferences or their affinity for different types of offers (see sidebar “What customer segmentation can look like”). With a more granular approach to customer segmentation, retailers can

## What customer segmentation can look like

**By grouping customers** into aggregate, anonymized categories, retailers can unlock additional margin. Here are groups to which marketers might offer different kinds of personalized promotions:

- ***Discount sensitive.*** Retailers could give steeper discounts to customers who might jump to a competitor if they don't receive a price reduction. A smaller margin is preferable to no sale at all.
- ***Product preferences.*** Consumer companies could deliver targeted promotions to groups that show affinity for particular products.
- ***Purchasing channel preferences.*** Retailers can offer discounts through specific channels, such as direct mail, in-app offers, or email.
- ***Infrequent buyers.*** For customers who have not made a purchase in some time, companies may send offers at a rate that's in line with past purchasing frequency.
- ***Loyalty program members.*** Marketers can serve more granular, personalized promotions to customers who have opted in to loyalty programs. The depth of the discount might depend on the membership level.

craft promotions that target specific customer life cycle stages (such as new-customer acquisition, customer retention, repeat purchase, or risk of churn) or specific business objectives (such as promoting a particular brand or category or encouraging cross-selling).

In a world saturated with promotions, companies can use targeted offers to stand out. Retailers that get this right can help ensure a better shopping experience while also enjoying better margins from saving on promotional costs and fueling more conversions. Ideally, marketers can develop a program of targeted offers at scale that accomplishes the following:

- applies business rules and algorithms to determine offerings and timings of delivery
- builds flexible, fit-for-purpose coupons (such as tiering discount rates so that those who buy more save more, limiting usage to certain categories or time periods, or designing offers that include or exclude certain premium categories or brands)
- delivers targeted promotions through all available marketing channels, such as through a company website or app, push notifications, text messages, or emails
- activates personalization with an always-on cadence where relevant
- accompanies offers with clear, highly relevant communications, such as with dynamic recommendations that update in real time for individual customers based on their purchasing or browsing history

With this variety of targeted offers, marketers can create a seamless omnichannel experience for customers, in which they receive targeted, streamlined promotions without conflicting or overwhelming information from other places. Companies should be smart about how much margin they're giving away when and where, encourage specific objectives rather than overly broad ones, and ensure that promotions are offered at the right time to the right people (see sidebar "How one retailer unlocked growth by launching targeted offers"). From what we've observed, companies that push incremental sales through targeted promotions can see a 1 to 2 percent lift in sales and a 1 to 3 percent improvement in margins.

## **Relevant marketing through gen AI-enhanced personalized content**

To truly enhance the impact of targeted promotions, companies can use gen AI to highly tailor copy and creative content that resonates more strongly with groups and subgroups of consumers than traditional marketing communications. Marketers have long aimed to improve the customer experience and influence consumer decisions by creating convenience or offering better prices. The next step is to make the buying experience even more convenient or enjoyable through greater relevance.

Traditionally, addressing small consumer groups with customized content has been cost-prohibitive and practically infeasible. Generative AI allows marketers to develop such content at scale at lower cost. While many marketers are currently piloting gen AI programs for this purpose, most are doing so manually with one-off experimental tools

# How one retailer unlocked growth by launching targeted offers

**One North American retailer** used a traditional calendar-based approach to promotions, offering mass discounts to all customers during holidays, and created a tiered discount program for different types of members of its loyalty program at other times.

Recently, the company embarked on a mission to pivot toward personalized and data-driven marketing. It wanted to know how to increase demand during the times of year with softer sales. The company explored how to deploy mass promotions in a personalized way—with offers to the full customer base during specific days, but with different depths of discounts to different customer segments. It also experimented with offers targeted only to specific customers who fit certain criteria.

To start, the marketing team developed analytical models to assess the likelihood that a customer

would respond positively to an offer—called “promotion propensity”—based on past purchases. The retailer then engaged in A/B testing for each model, using data analytics that allowed agile, cross-functioning teams to test the value of various targeted offers (made primarily via email) over a series of two-week sprints. After learning that its customers felt overwhelmed by too many promotions, the company pared back some of its offer frequency and simplified the customer experience.

After three months of using this more targeted approach, the retailer saw a boost of about 3 percent in annualized margins during initial tests. The company now plans to expand targeted promotions at scale.

and use cases; they are not automating or integrating to reduce the bottlenecks of operational inefficiencies, nor are they evaluating content performance.

Smarter use of gen AI can help unlock more cohesive personalization opportunities with touchpoints and interactions more tailored to what customers want. Currently, there is no consolidated suite that does this end to end. But with thoughtfully integrated solutions that align people, processes, and platforms, marketers can deliver more rigorous content creation, smoother orchestration across teams, and seamless implementation of targeted promotions and content.

While content creation today is highly manual, gen AI can accelerate and magnify the entire process, helping channel operators, creators, and analysts become more productive. As more marketing material gets fed into a robust, foundational content data model, gen AI can learn from a feedback loop and create more copy that can be tailored for personalization. One important thing to note: As organizations increasingly use gen AI, it is critical that they build models to validate and govern gen AI—created content in order to establish guardrails against bias, toxicity, and hallucinations, and to ensure that content is in accordance with enterprise standards and design systems.

To see what's possible for better workflow in the future, it helps to look at how the ecosystem operates—both today and in the future—across the three stages of content production (table).

Table

	Content creation today	Content creation with gen AI
Production and versioning	Channel operators (who have ownership over specific channels) typically design campaigns and creative briefs with clear parameters. Content creators then usually conduct research, hold ideation and iteration sessions with creative teams, and produce versions of content to use. They tag these text and imagery assets with metadata so they can be discoverable and store them in a digital-asset-management (DAM) system for future retrieval.	Using natural-language processing, gen AI can help content creators brainstorm new ideas and concepts. It can also assist with writing text and selecting what formats (such as static imagery or videos) specific audiences might prefer to see. Content creators can also use gen AI to identify already existing creative assets and for versioning final assets for all channel placements. Gen AI can tag final assets with metadata for storage in a DAM system, improving efficiency for finding assets and decreasing reliance on content creators following tagging standards.
Activation	Channel operators often create business rules to determine the best course of action based on customer data and hypothetical interactions. They populate a content data model with creative assets and use decisioning services to select the best text and images. Using the tags from content creators, they then retrieve the correct creative files from the DAM and feed them into the appropriate channel platform according to campaign requirements. Based on the content data model, they then publish the content, sending the appropriate creative files to relevant audiences.	Gen AI tools can process requests to retrieve content from a DAM system and deliver creative assets with the right file sizes, resolution, and formats to fit within guidelines for specific channels. They can also assign the appropriate campaign tags by channel to ensure robust campaign metadata for measurement.
Performance and measurement	Marketing analysts measure the performance of various campaigns, provide and categorize insights, and develop dashboards that offer windows into customer experience, channel performance, and content performance. Throughout the content creation process, data scientists develop propensity models that measure the likelihood that a customer will take action, whether it's at the top of the marketing funnel (such as viewing or clicking on a piece of content) or closer to the bottom of the funnel (when a customer makes a purchase).	With gen AI, marketers can develop a standardized measurement approach, using content metadata from a DAM system (such as size, color, and theme of the content), campaign metadata, campaign performance data, and decisioning services to build better performance tracking of content and campaigns. Marketing leadership can better oversee performance by creative version, audience, channel, or campaigns.

Our experience so far shows speedy results. We've seen some marketers deploy gen AI to personalize content development 50 times faster than a more manual approach (see sidebar “How a European telecom used gen AI to enhance marketing materials”).

## Technology as a foundational differentiator

To better target promotions and gen AI–boosted content, companies can turn to a tech stack that brings everything together. In 2019, McKinsey published “A technology blueprint for personalization at scale,” which described a “4D” strategy (data, decisioning, design, and distribution) for marketing technology.<sup>2</sup> We add one more critical element to this: measurement (exhibit). (It's not a “D,” but it's just as relevant.) Marketers can establish a solid foundation for growth through personalization by ensuring these five elements use the latest technological innovations and integrate with each other seamlessly (see sidebar “A tech-enabled evolution from mass discounts to targeted offers”).

### Data

By improving data collection and analysis, marketers can gain deeper insights into customer behaviors and preferences. And while many enterprises have invested in data lakes (storage platforms that hold, process, and analyze structured and unstructured data) and customer data platforms (software that centralizes and unifies customer data from multiple sources to create a single view of each customer), better targeted offers and content requires expanding data architecture in five categories:

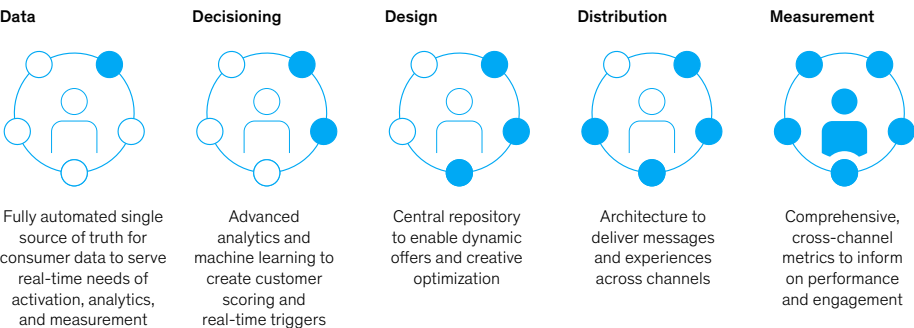
- a promotions subject area that includes the history of offers and redemptions
- a content subject area that includes the history of content delivery and engagement
- universal (and potentially gen AI–enabled) metadata and taxonomy, which can improve the flow of automation

<sup>2</sup>Sean Flavin and Jason Heller, “A technology blueprint for personalization at scale,” McKinsey, May 20, 2019.

### Exhibit

## Marketers can establish a solid foundation for personalization with an effective technology framework.

### Technology blueprint for personalization at scale



# How a European telecom used gen AI to enhance marketing materials

**European telecom company** recently boosted its marketing strategy by integrating a personalization engine that uses both AI and gen AI. The telco previously relied on mass promotions and a calendar-based approach. To better engage its customers, it set up a next-best-action engine, using a combination of multiple machine learning models to determine the most effective actions to suggest to each customer.

The next-best-action mechanism predicted the probability of a customer accepting a specific action and the expected value if the offer was accepted. It then ranked these actions, optimizing for the highest expected value from a response. The company then tested a granular set of about 2,000 different actions by texting messages to customers.

The marketing team then deployed gen AI-enhanced messaging for a handful of different

campaigns to see how a more personalized approach would work. It sent messages to customers based on factors such as age, gender, and data usage. The copy blended general company messaging with campaign-specific features, making the offers feel more personal and relevant to customers. Guardrails were put in place to limit the length, tone, and content, ensuring that messages were concise, relevant, and respectful of user privacy.

Over the course of a few months, the experiment showed that customers receiving the personalized messages from the gen AI-enhanced campaigns engaged and took action 10 percent more often than customers who did not receive personalized content. The telecom is now looking at more extensive use of personalized content across multiple marketing channels.

- a robust analytics infrastructure and MLOps (machine learning operations), with dedicated feature stores to rapidly deploy and scale ML models
- new data pipelines and integrators, prompt stores, and vector databases to build or customize large language model implementations

Along with granular transaction data, these additional data assets are the backbone of AI-powered decisioning to predict customer behavior in any channel and can be a true differentiator in today's market.

## Decisioning

To develop new targeted promotions and content through more robust targeting, companies can also benefit from refreshing their decision engines with new AI models. Their tasks include the following:

- **Promo propensity** predicts the likelihood of a customer making a purchase due to a promotion, based on previous customer purchasing and engagement behavior. This can lead to better customer satisfaction and targeting with the right level of discount to improve margins.

# A tech-enabled evolution from mass discounts to targeted offers

**A large, established North American retailer,** known for its deep discounts during sales, transformed itself to enhance pricing and promotional strategies. A few years ago, the company began to develop an incremental marketing approach to improve profit. This shift aimed to optimize the financial outcomes of marketing efforts, ensuring that each promotion was both attractive and effective.

To achieve this, the company incorporated three key levers: technology, analytics, and activation. The technology team integrated the company's legacy point-of-sale (POS) infrastructure with its marketing technology stack, creating use cases that spanned both systems. This integration allowed for seamless data sharing and a unified

view of customer behavior. On the analytics side, the retailer built models that provided unique insights into the overlap between products and customers, which helped it tailor its offers more precisely and prioritize retention efforts for frequent shoppers who had stopped purchasing. To bring it all together, the company formed cross-functional teams with representation from marketing, pricing, technology, and operations to align all commercial stakeholders and launch targeted offers while scaling back on mass-market promotions.

The results were dramatic. Over a single year, the company produced \$400 million in value from initial pricing improvements, and an additional \$150 million from gen AI-enabled targeted offers.

- **Promo uplift** (or promo effectiveness) predicts promotion ROI (uplift) by analyzing customer behavior during promotion and no-promotion periods.
- **Content propensity** predicts the likelihood that a customer will respond to a piece of content. Propensity scores enable the automated delivery of the best content that will encourage a customer to respond to a call to action.
- **Content effectiveness** measures the effectiveness of content by analyzing customer response. Highly effective content can be reused or thematically replicated in future campaigns.

Model outputs are then fed to a decision engine that ranks and determines the best offer and content to show a customer at a given point in time.

## Design

Innovative design ensures that content is both engaging and relevant. A sophisticated design layer that oversees two critical workflows (offer management and content production) helps manage the process, fueling both operational excellence and agility.



Targeted offers work best with an integrated offer management system to catalog, manage, deliver, and redeem them across any channel, including e-commerce and point of sale. Meanwhile, content management begins with gen AI tools for creating copy and developing creative assets, handling versioning, and auto-formatting content as varied as billboards and mobile devices. Digital assets are stored in a single, centralized digital-asset-management (DAM) system.

It's crucial that both the offer management and DAM systems are well integrated into all downstream channels. This enables easy search, reuse, and dynamic delivery of assets.

### **Distribution**

Achieving true, real-time personalization requires sophisticated architecture that delivers seamless and consistent messaging to the right audiences at the right time as customers traverse channels.

This critical infrastructure combines core capabilities:

- instant processing of customer signals, fed to journey orchestration and decisioning platforms to optimize the right message and channel across customer touchpoints
- front-end tools that support a company's website, app, and email (such as content and campaign management systems, as well as dynamic content optimization), built with dynamic modular templates and API integrations to render personalized content in real time
- interoperability and integrations across multivendor platforms

### **Measurement**

A comprehensive marketing technology stack requires thorough measurement to facilitate ongoing optimization and improvement. To validate the ROI of personalization efforts, rigorous incrementality testing, standardized performance metrics, and measurement playbooks are essential.

Businesses need actionable intelligence for continuous improvement. Marketers can implement closed-loop measurement by aggregating data from all channels into a centralized reporting engine that produces self-serve dashboards for distinct stakeholders—from executive leadership tracking revenue and margin impact to marketing operators optimizing campaigns in real time.

---

To unlock the next frontier of gen AI-enabled personalization, marketers can begin with a thorough assessment of their opportunities by making the following moves:

- mapping out the areas where targeted offers and more relevant content can drive the highest value
- identifying the lifetime value events that they want to encourage

- conducting a tech diagnostic to identify missing tools
- reconfiguring processes in areas such as talent, data, tech, analytics, and marketing operating models to optimize for targeted promotions and content development

Excellence in execution sets the leaders apart from the followers. Success hinges on seamlessly integrated platforms backed by well-trained teams that can fully maximize investments. To improve performance, marketers can focus on operational efficiency, eliminate redundant systems, and establish robust governance. This can help bring together disparate tools into a unified engine for more relevant and personalized customer engagement, contributing to real growth. [Q](#)



**Eli Stein** is a partner in McKinsey's Bay Area office; **Kelsey Robinson** is a senior partner in the Boston office; **Alexis Wolfer** is an associate partner in the Southern California office; **Gaelyn Almeida** is an expert associate partner in the Washington, DC, office; and **Willow Huang** is an associate partner in the New York office.

The authors wish to thank Carolyn Spalding, Jianing Cheng, and Molly McCulloch for their contributions to this article.

Copyright © 2025 McKinsey & Company. All rights reserved.